

Example 1: Creating a table using HBase Shell

HBase has a command-line shell

```
hbase(main):001:0>
```

This shell can be used to **create tables, insert and read data etc**

```
hbase(main):001:0> create 'notifications','attributes','metrics'
```

Create a table to hold
notifications data

```
hbase(main):001:0> create 'notifications', 'attributes', 'metrics'
```

This is the
name of the
table

```
hbase(main):001:0> create 'notifications' 'attributes', 'metrics'
```

These are column families

column families

```
hbase(main):001:0> create 'notifications' 'attributes', 'metrics'
```

Column families are groups
of columns which are
usually semantically related

column families

```
hbase(main):001:0> create 'notifications' 'attributes', 'metrics'
```

When you create a table in HBase, you don't have to specify the columns in the table

column families

```
hbase(main):001:0> create 'notifications' 'attributes', 'metrics'
```

For instance, this notification table might have columns like **type**, **text**, **timestamp** etc

column families

```
hbase(main):001:0> create 'notifications' 'attributes', 'metrics'
```

columns like type, text, timestamp etc

These columns are defined on the fly when you insert data for a specific a row id

column families

```
hbase(main):001:0> create 'notifications' 'attributes', 'metrics'
```

Every column has to
belong to some
column family

column families

```
hbase(main):001:0> create 'notifications','attributes','metrics'
```

The attributes column family may have columns like **type**, **timestamp**

column families

```
hbase(main):001:0> create 'notifications', 'attributes', 'metrics'
```

The metrics column family may have columns like #clicks, # views

column families

```
hbase(main):001:0> create 'notifications','attributes','metrics'
```

Every table must
have at least 1
column family

column families

```
hbase(main):001:0> create 'notifications','attributes','metrics'
```

Column families, unlike columns are usually created at the time of table creation

column families

```
hbase(main):001:0> create 'notifications','attributes','metrics'
```

It is possible to add or
change column families
later, but this is rarely done

Example 2: Inserting data into a table using HBase Shell

put

```
put 'notifications',2, 'attributes:for_user','Chaz'
```

Data is inserted into
HBase tables using
the put operation

```
put 'notifications',2, 'attributes:for_user','Chaz'
```

id	type	for user	from user	timestamp
1	Friend request status	Ryan	Jessica	146710201
2	Comment	Chaz	Daniel	146711200
3	Comment	Rick	Brendan	1467112205
4	Like	Rick	Brendan	1467112213

Each entry in a HBase table is like a cell in a traditional table

row	column	value
2	for user	Chaz

row	column	value
2	for user	Chaz

```
put 'notifications', 2, 'attributes:for_user', 'Chaz'
```

With `put`, you insert data
1 cell at a time

row	column	value
2	for user	Chaz

```
put 'notifications', 2, 'attributes:for_user', 'Chaz'
```

The table name

row	column	value
2	for user	Chaz

```
put 'notifications', 2, 'attributes:for_user', 'Chaz'
```

id	type	for user	from user	timestamp
1	Friend request status	Ryan	Jessica	146710201
2	Comment	Chaz	Daniel	146711200
3	Comment	Rick	Brendan	1467112205
4	Like	Rick	Brendan	1467112213

The row id

row	column	value
2	for user	Chaz

```
put 'notifications', 2, 'attributes:for_user', 'Chaz'
```

The column

id	type	for user	from user	timestamp
1	Friend request status	Ryan	Jessica	146710201
2	Comment	Chaz	Daniel	146711200
3	Comment	Rick	Brendan	1467112205
4	Like	Rick	Brendan	1467112213

row	column	value
2	for user	Chaz

```
put 'notifications', 2, 'attributes:for_user', 'Chaz'
```

id	type	for user	from user	timestamp
1	Friend request status	Ryan	Jessica	146710201
2	Comment	Chaz	Daniel	146711200
3	Comment	Rick	Brendan	1467112205
4	Like	Rick	Brendan	1467112213

The value

row	column	value
2	for user	Chaz

```
put 'notifications',2, 'attributes:for_user','Chaz'
```

The column name is
specified along
with it's **column
family**

row	column	value
2	for user	Chaz

```
put 'notifications',2, 'attributes:for_user','Chaz'
```

Every column
must belong to a
column family

row	column	value
2	for user	Chaz

```
put 'notifications',2, 'attributes:for_user','Chaz'
```

An HBase table is
like a sorted map

Key

row	column	value
2	for user	Chaz

```
put 'notifications', 2, 'attributes:for_user', 'Chaz'
```

Key Value

row	column	value
2	for user	Chaz

```
put 'notifications',2, 'attributes:for_user', 'Chaz'
```

Key	Value
put 'notifications', 2, 'attributes:for_user', 'Chaz'	
put 'notifications', 2, 'attributes:type', 'Comment'	

With the put operation, we are
inserting new keys into the map

Example 3: Updating data using the HBase Shell

put (contd)

The put operation has **2 purposes**:

Inserting values for new keys (**row id, column**)

Updating the value for **existing keys**

Inserting values for new keys

Say we want to **track the number of times** a notification was opened/clicked on

```
put 'notifications',2, 'metrics:open',0
```

Updating the value for existing keys

Inserting values for new keys

```
put 'notifications', 2, 'metrics:open', 0
```

Create row id 2 in notifications table
If the row id does not exist

Updating the value for existing keys

Inserting values for new keys

```
put 'notifications', 2, 'metrics:open', 0
```

Create column **open** for the row id 2 in the
metrics column family

Updating the value for existing keys

Inserting values for new keys

```
put 'notifications', 2, 'metrics:open', 0
```

Insert the value 0 for the **row id 2** and
column open

Updating the value for existing keys

Updating the value for existing keys

```
put 'notifications', 2, 'metrics:open', 0
```

```
put 'notifications', 2, 'metrics:open', 1
```

When someone clicks or opens the notification, update the value for this key

Updating the value for existing keys

```
put 'notifications',2, 'metrics:open',0  
:  
put 'notifications',2, 'metrics:open' 1
```

Since the row id and column already exist, the value is just updated

Updating the value for existing keys

```
put 'notifications',2, 'metrics:open' 0  
:  
put 'notifications',2, 'metrics:open' 1
```

HBase actually does not lose
the old value

Updating the value for existing keys

```
put 'notifications',2, 'metrics:open',0  
.  
put 'notifications',2, 'metrics:open',1
```

The history of updates for
a specific key is maintained
and retrievable

Updating the value for existing keys

```
put 'notifications',2, 'metrics:open',0 1467184097569  
put 'notifications',2, 'metrics:open',1 1467181276487
```

Each version is stored with
the created timestamp

Updating the value for existing keys

```
put 'notifications',2, 'metrics:open',0 1467184097569  
put 'notifications',2, 'metrics:open',1 1467181276487
```

During retrieval, the latest version is retrieved by default

Example 4: Retrieving data using the HBase Shell

get

```
get 'notifications',2
```

Data is retrieved from
HBase tables using the
get operation

```
get 'notifications',2
```

You can retrieve
data for 1 row id
at a time

```
get 'notifications', 2
```

The default behavior is
to return all the columns
for the specified row id

```
get 'notifications', 2
```

get has **2** mandatory
arguments

```
get 'notifications', 2
```

The table name

```
get 'notifications', 2
```

The row id

```
get 'notifications',2
```

COLUMN

```
attributes:for_user  
attributes:type  
metrics:open
```

CELL

```
timestamp=1467179977249, value=Chaz  
timestamp=1467181276487, value=Comment  
timestamp=1467184097569, value=1
```

The column names with
their column families

```
get 'notifications',2
```

COLUMN

attributes:for_user
attributes:type
metrics:open

CELL

timestamp=1467179977249, value=Chaz
timestamp=1467181276487, value=Comment
timestamp=1467184097569, value=1

The values for these
columns

```
get 'notifications',2
```

COLUMN

attributes:for_user
attributes:type
metrics:open

CELL

timestamp=1467179977249	value=Chaz
timestamp=1467181276487	value=Comment
timestamp=1467184097569	value=1

The timestamp when the value was last updated

```
get 'notifications',2
```

COLUMN

attributes:for_user
attributes:type
metrics:open

CELL

timestamp=1467179977249, value=Chaz
timestamp=1467181276487, value=Comment
timestamp=1467184097569, value=1

In HBase, **value + timestamp**
is called a **cell**

```
get 'notifications',2
```

COLUMN

attributes:for_user
attributes:type
metrics:open

CELL

timestamp=1467179977249, value=Chaz
timestamp=1467181276487, value=Comment
timestamp=1467184097569, value=1

A row id, column can have multiple cells

```
get 'notifications',2
```

COLUMN

attributes:for_user
attributes:type
metrics:open

CELL

timestamp=1467179977249, value=Chaz
timestamp=1467181276487, value=Comment
timestamp=1467184097569, value=1

By default, the cell with the latest timestamp is retrieved

```
get 'notifications',2,'metrics:open'
```

COLUMN
metrics:open

CELL
timestamp=1467184097569, value=1

You can ask get to
retrieve values for
specific columns

```
get 'notifications', 2, 'metrics:open', 'attributes:type'
```

COLUMN

attributes:type
metrics:open

CELL

timestamp=1467181276487, value=Comment
timestamp=1467184097569, value=1

You can also specify
a list of columns

Example 5: Retrieving a range of row ids using the HBase Shell

scan

get operations only
allow you to retrieve
1 row id at a time

HBase tables are sorted
maps, ie. row ids are sorted

With the scan operation,
you can retrieve row ids
within a specified range

```
scan 'notifications'
```

ROW	COLUMN+CELL
1	column=attributes:text, timestamp=1467179437142, value="Hi there! Buy this thing "
1	column=attributes:type, timestamp=1467179418482, value=promotion
2	column=attributes:for_user, timestamp=1467179977249, value=Chaz
2	column=attributes:type, timestamp=1467181276487, value=Comment
2	column=metrics:open, timestamp=1467184097569, value=1

Returns all values from
the notifications table

```
scan 'notifications'
```

ROW	COLUMN+CELL
1	column=attributes:text, timestamp=1467179437142, value="Hi there! Buy this thing "
1	column=attributes:type, timestamp=1467179418482, value=promotion
2	column=attributes:for_user, timestamp=1467179977249, value=Chaz
2	column=attributes:type, timestamp=1467181276487, value=Comment
2	column=metrics:open, timestamp=1467184097569, value=1

All columns for row id 1

```
scan 'notifications'
```

ROW	COLUMN+CELL
1	column=attributes:text, timestamp=1467179437142, value="Hi there! Buy this thing "
1	column=attributes:type, timestamp=1467179418482, value=promotion
2	column=attributes:for_user, timestamp=1467179977249, value=Chaz
2	column=attributes:type, timestamp=1467181276487, value=Comment
2	column=metrics:open, timestamp=1467184097569, value=1

All columns for row id 2

```
scan 'notifications',{COLUMNS => ['attributes:type'],LIMIT => 1,STARTROW => "2"}
```

You can pass in a
dictionary with some
specifications to scan

```
scan 'notifications',{COLUMNS => ['attributes:type'],LIMIT => 1,STARTROW => "2"}
```

COLUMNS

A list of column names
(with column family)

```
scan 'notifications',{COLUMNS => ['attributes:type'],LIMIT => 1,STARTROW => "2"}
```

LIMIT

The number of values to
be returned

```
scan 'notifications',{COLUMNS => ['attributes:type'],LIMIT => 1,STARTROW => "2"]
```

STARTROW

Only values starting
from this row id

You can specify the end
row id using STOPROW

```
scan 'notifications',{COLUMNS => ['attributes:type'],LIMIT => 1,STARTROW => "2"]
```

STARTROW

STOPROW

Note: The row id has to be passed as a string

```
scan 'notifications' {COLUMNS => ['attributes:type'],LIMIT => 1,STARTROW => "2"}
```

COLUMNS
LIMIT
STARTROW
STOPROW

Any or all of these options can be specified

Example 6: Deleting data using the HBase Shell

delete

```
delete 'notifications',2, 'attributes:for_user'
```

Data is deleted from
HBase tables using
the delete operation

row id	column
2	for user

```
delete 'notifications', 2, 'attributes:for_user'
```

Delete data 1 cell at a time

row id	column
2	for user

```
delete 'notifications',2, 'attributes:for_user'  
get 'notifications',2
```

Before delete

COLUMN	CELL
attributes:for_user	timestamp=1467179977249, value=Chaz
attributes:type	timestamp=1467181276487, value=Comment
metrics:open	timestamp=1467184097569, value=1

This cell is deleted

```
get 'notifications',2
```

Before delete

COLUMN

 attributes:for_user
 attributes:type
 metrics:open

CELL

 timestamp=1467179977249, value=Chaz
 timestamp=1467181276487, value=Comment
 timestamp=1467184097569, value=1

After delete

COLUMN

 attributes:type
 metrics:open

CELL

 timestamp=1467181276487, value=Comment
 timestamp=1467184097569, value=1

Example 7: Deleting a table using the HBase Shell

drop

Delete an HBase table using the drop command

```
drop 'notifications'
```

```
ERROR: Table notifications is enabled. Disable it first.
```

Before deleting a table,
it must be disabled first

```
drop 'notifications'
```

ERROR: Table notifications is **enabled** Disable it first.

When a table is in use,
HBase keeps an index of
the row ids in memory

```
drop 'notifications'
```

ERROR: Table notifications is **enabled** Disable it first.

HBase also keeps a log of recent changes in memory, which are periodically flushed to disk

```
drop 'notifications'
```

ERROR: Table notifications is enabled. **Disable** it first.

Disabling will flush all recent changes to disk and remove the row id index from memory

```
disable 'notifications'
```

```
drop 'notifications'
```

```
list
```

Use the **list** command to check whether the table has been dropped

```
disable 'notifications'
```

```
drop 'notifications'
```

```
list
```

Before drop

```
TABLE  
notifications  
test
```

After drop

```
TABLE  
test
```